



Multi-Instance Codex CLI for Safe, Scalable Development

A multi-instance (multi-agent) setup can accelerate development by running specialized Codex processes in parallel (e.g. a *scaffold* agent, a *tester* agent, etc.) while isolating their access. This approach brings benefits like role separation and throughput, but also introduces complexity and risk. In the context of a modular-monolith (FastAPI + Next.js + PostgreSQL + Docker) for a small team, we explore **when and how** to safely run multiple Codex CLI instances. We cover use-cases, feature mapping, orchestration patterns (local and CI), security, audit/provenance, governance, migration rollback, and provide concrete templates (config, YAML, rules, prompts, scripts) and checklists. Recommendations are grounded in official Codex CLI docs and best practices (OWASP, GitHub Actions, SRE/CICD guidance).

Executive Summary (top recommendations) – Use a single Codex instance by default. Scale to multiple *scaffold*, *fixer*, *reviewer* agents only when justified (see decision matrix). Enforce **least privilege** by default (read-only sandbox), require human review of generated code (never auto-merge), and pin all Codex versions/prompts for reproducibility. In CI (GitHub Actions), run Codex steps in isolated jobs with per-job secrets. Always log every Codex run (JSONL output, prompts, versions) to an audit store. Provide `execpolicy` rules to block dangerous commands (e.g. destructive shell ops). Use skills (via `.agents/skills`) and AGENTS.md to encapsulate domain knowledge. Prioritise immediate actions: set up CODEX_HOME, create `AGENTS.md`, configure CI jobs for lint/test, draft basic `codex` rules, and test one Codex “scaffold” run in a feature branch.

1. Rationale & Use-Cases for Multi-Instance Codex

For a 1-3 developer team, a **single Codex CLI instance** (per task) is simplest and often sufficient. However, certain scenarios motivate multiple instances (see Table 1):

- **Role separation:** Run one agent focused on scaffolding code, another on writing tests, another on docs/reviews. Each agent uses a tailored instruction set, reducing prompt complexity.
- **Parallel tasks:** Perform code generation and concurrent checks (lint, security analysis) simultaneously, speeding up feedback cycles.
- **Isolation & security:** Limit each instance’s privileges to a subset of project (e.g. front-end vs back-end directories).
- **Scaling automation:** In CI/CD, use multiple agents to process PRs, security, and release steps in parallel.

For a small project, multi-instance is **optional**. Use it when project size or pipeline complexity grows. Table 2 (later) offers signals: e.g. repeated long Codex runs, distinct code domains (finance vs dev modules), or backlog of manual tasks. As GitHub issue #3280 noted: “*We should be able to orchestrate coding agents... since each one has a dedicated domain*” ¹. In practice, teams have run “*multi-agent security pipelines that catch SQL injection before code review*” ², demonstrating real-world gains.

Table 1. Use-cases for multiple Codex instances vs single-instance trade-offs.

Scenario	Single Instance	Multiple Instances
Simple CRUD tasks	Easy setup, single prompt context.	Overkill; additional complexity.
Parallel operations	Must run tasks one-by-one (slower).	Can run simultaneously (faster feedback ³). But be wary: parallel heavy tasks can crash Codex ³ .
Domain separation	One large prompt with mixed context.	Agents per domain (front-end vs back-end) – each uses relevant AGENTS.md guidance ⁴ .
Security checking	Codex or external tools in pipeline.	Dedicated <i>security agent</i> can pre-scan code changes (e.g. SQLi detection as in Rezvani's example ²).
Maintenance speed	Simpler, lower overhead.	High initial setup/coordination cost, but scales tasks.
Risk management	Fewer moving parts, easier auditing.	Need strict sandbox rules and logs per agent.

In summary, multi-instance is a **scaling** technique: use selectively. For a 1–3 person team, start single-instance (monolith-first), then evolve to parallel agents when productivity bottlenecks (e.g. code review backlog, long refactors) appear.

2. Codex CLI Features Mapped to Multi-Instance Patterns

We align Codex CLI capabilities with multi-agent workflows. Where official docs lack specifics, we note it as “*unspecified*” and suggest conservative defaults.

Interactive vs Non-Interactive

- **Interactive (TUI) mode** (just `codex`) is for exploratory, one-off sessions. Not suitable for automation in CI because it requires a terminal. Multi-instance use in CI requires **non-interactive mode** (`codex exec`) ⁵.
- `codex exec` is scriptable. It accepts a prompt (inline or via `-`), runs the plan/tools, and can output JSONL (via `--json`) ⁶. This is essential for CI and parallel runs. Each `codex exec` is one session; to run multiple, simply launch multiple `codex exec` processes.

Sandboxing Modes

Codex supports three sandbox levels: `read-only`, `workspace-write`, `danger-full-access` ⁷. For multi-instance:

- **Default to `read-only`**: safest (no file writes) for any agent that should not modify code, e.g. reviewers or linters. This enforces safety by default (low effort, high long-term benefit).
- `workspace-write`: allows editing files. Use only for scaffold/implement agents. Combine with *approval* settings to gate changes.
- `danger-full-access`: unrestricted. **Avoid** unless absolutely needed (e.g. massive refactor in single dev context) as it breaks containment. We treat it as prohibited for multi-agent CI.

Use **profiles** in config.toml to pre-set sandbox per role. Example `.codex/config.toml` (later) will show roles (dev, test, doc) with sandbox and approver settings.

AGENTS.md and Rules

- **AGENTS.md** provides global/project instruction chain ⁸. For multi-instance, each agent shares the same repository, but we can use **per-agent overrides** via CODEX_HOME or working directory tricks:
 - Set `CODEX_HOME` or `--profile` per instance to isolate global AGENTS.md if needed (e.g. one agent only sees global rules, another only sees project-specific).
 - Each agent can be launched from a subdirectory with its own `AGENTS.override.md` to specialize instructions for that role (see [24tL295-L303]).
- **Rules and ExecPolicy:** Use `.codex/rules/*.yaml` files with allow/deny blocks. This is crucial for multi-agent safety. For example, a CI *fixer* agent should *not* delete or create PRs, so forbid `rm -rf` or `git push`. Use `codex execpolicy check` to validate rules ⁹. (See templates below.)

JSONL, Version Pinning, Auth

- **JSONL output:** In CI, always run with `--json` to record every step (prompts, actions, message deltas) as newline-delimited JSON. This stream is key for audit logs. Enable `--output-schema` if expecting structured data from Codex (like JSON-based API spec).
- **Session Resume:** In CI runs, keep sessions short; no need to resume across jobs. In local dev, `codex exec resume --last` can continue a paused session, but for reproducibility avoid long-lived sessions in CI.
- **Auth modes:** In CI, use an API key with limited scope (e.g. ChatGPT API key or GPT-5 Codex key). **Unsure** if Codex CLI can use ephemeral tokens from OIDC; likely not. For now, store API key in GitHub Secrets (read-only for runners) and mount to `CODEX_API_KEY` env var. On local, either ChatGPT account login (with quotas) or API key. Document both options.
- **Version pinning:** Always pin Codex CLI version (e.g. via `brew install codex@x.y` or Docker image). Also pin model (e.g. `--model gpt-5.3-codex`). Save the CLI binary in Docker image or via GitHub Action checkout. Reference [20tL343-L352] for model usage rates.

3. Orchestration Patterns

Local Development

Per-Instance Settings

Each Codex process should use its own home/config to avoid state collision. Strategies:
- Set `CODEX_HOME` per instance (e.g. `CODEX_HOME=$PWD/.codex-sandbox`). This isolates history, chats, and skills.
- Use `--profile` for different roles. Example: in `.codex/config.toml` define profiles `scaffold`, `fix`, `review`, each with different `sandbox` and `credential`:

```
[profile.scaffold]
sandbox = "workspace-write"
ask_approval = "on-change"
[profile.fixer]
```

```
sandbox = "read-only"
ask_approval = "never"
[profile.reviewer]
sandbox = "read-only"
ask_approval = "on-request"
```

(Expand later in appendix.)

Launching Multiple Agents

For example, a developer on branch `feature/foos` might run:

```
# Scaffold agent (writes code)
CODEX_HOME=./codex_scaffold codex exec --profile scaffold "Create boilerplate
for feature 'foo'"
# Test-fixer agent (non-writing, reviews code)
CODEX_HOME=./codex_fixer codex exec --profile fixer "Review code in ./apps/
api for security issues"
```

These run in parallel shell terminals. Each has separate CODEX_HOME so they don't share context unless explicitly allowed. If both open windows, they won't conflict.

Per-Project / Per-Directory Skills

Place skills under the repo: e.g. `.agents/skills/{scaffold, fixer, reviewer}/SKILL.md`. Codex loads them by name or via implicit triggers ¹⁰. For instance, a `scaffold` skill might contain templates or generators. Skills reduce repeated prompting overhead.

CI (GitHub Actions)

Use multiple jobs (in a workflow) to run Codex tasks in parallel:

- **Job matrix:** define roles. Example `matrix: [scaffold, fixer, reviewer]` each job sets `CODEX_HOME: ${{ github.workspace }}/codex_$role`.
- **Secrets per job:** restrict secrets to needed scopes. For instance, only the fixer job might have a DB read secret if running DB checks, while scaffold only needs GitHub token.
- **Ephemeral runners:** each job runs on a clean VM/container, isolating instances completely.
- **Gating and PR-only:** require runs on PR branches, not main. Disable write access (`allow-bots: false` or manual review required).
- **Artifacts:** upload JSON logs (`--json > log.json`) and Codex `final_message.txt` to GitHub as artifacts for audit.

Example GitHub Actions snippet (simplified):

```
name: Codex Agents
on: [pull_request]
jobs:
  codex:
```

```

strategy:
  matrix:
    role: [scaffold, fixer, reviewer]
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v3
    - name: Setup Codex CLI
      uses: openai/setup-codex@v1
      with:
        version: 5.3
    - name: Run Codex agent
      env:
        CODEX_HOME: ${{ github.workspace }}/codex_${{ matrix.role }}
        CODEx_API_KEY: ${{ secrets.CODEX_${{ matrix.role }}_KEY }}
      run: |
        case "${{ matrix.role }}" in
          scaffold) codex exec --json --profile scaffold ... ;;
          fixer)   codex exec --json --profile fixer ... ;;
          reviewer) codex exec --json --profile reviewer ... ;;
        esac
    - name: Upload Codex logs
      uses: actions/upload-artifact@v3
      with:
        name: log-${{ matrix.role }}
        path: codex_${{ matrix.role }}/log.json

```

This job runs all three roles in parallel (one per runner). Adjust prompts and secret usage per role.

(Optional) GitLab CI snippet would be similar using `parallel` matrix in `.gitlab-ci.yml`.

Hybrid (Local + CI + Runners)

For a mix of local and CI usage: ensure consistent setups. Use the same `codex/config.toml` and `AGENTS.md` in repo. Encode any CI-only settings (like `--output-file`) in the workflow, not in code. Developers can replicate CI runs locally via a `make` target.

4. Security & Secrets

Finance apps demand high security. We recommend **conservative defaults**:

- **Secrets Storage:** Never hard-code API keys. In GitHub Actions, store keys as *repository secrets* (encrypted by GitHub). For local dev, consider using environment variables or a secrets manager (e.g. AWS Secrets Manager, Vault). OWASP advises centralizing and rotating secrets ¹¹.
- **Least Privilege:** Always restrict Codex credentials to minimal scope. For API keys, grant only Codex access. For GitHub tokens, set default permissions to read-only unless needed ¹². Do not give Codex agent a user account with broad privileges.
- **Network Isolation:** In CI, run Codex on GitHub's runners; it cannot reach internal DB or cloud APIs unless secrets allow. For added safety, use OIDC to fetch temp creds if supported (currently unspecified).

- **Sandboxing & Rules:** Use Codex's sandboxing plus your own rules. Forbid dangerous operations in rules: e.g. disallow `sudo`, `kill`, `rm -rf /*`, `curl http://`, etc. Provide an example `rules.yaml` (below) that sets most commands to `prompt` or `deny`. The `execpolicy` command can verify these rules.
- **Prompt Injection Mitigation:** Ensure user inputs (like issue content or PR text) are not blindly fed into Codex prompts. Always escape or parameterize untrusted data, and avoid including raw user content in `AGENTS.md`. Codex CLI does not currently auto-sanitize prompts, so treat it like any interpreter: quotes, escapes, or use `--no-color` / `--output-schema` to validate.
- **Audit Logging:** Log every invocation (prompts and actions). GitHub Actions automatically record logs, but also **explicitly save the JSONL** (using `--json`). Structured logs should include at least: timestamp, agent role, CLI version, model, session ID, prompt hash, commit SHA. Store these as artifacts or push to a log store (or both).

Secrets Guidelines (OWASP/GitHub)

- *Store secrets in platform-provided vaults*: e.g. GitHub Secrets. Do not put in code or plaintext files ¹³.
- *Rotate frequently*: Use short-lived tokens if possible. OWASP recommends rotation and pipeline automation ¹⁴.
- *Audit usage*: Regularly check who has access to secrets; use GitHub audit logs.
- *Never echo secrets*: Use GitHub's mask or `::add-mask::` if printing is needed ¹⁵.

5. Auditability & Provenance

Each Codex run produces an event trail. We must capture and store this for auditing and reproducibility.

Logging Codex Invocations

- **JSONL Output:** Always run with `--json` to get newline-delimited JSON events ⁶. Save this (e.g. `codex.log`). Upload it as CI artifact (for PR runs) and keep locally after dev runs.
- **Metadata:** Alongside JSON, record: CLI version (`codex version`), model name, `AGENTS.md` hash, rules file hash, Git commit SHA, runner (hostname/ID), and start timestamp. This can be one JSON record or YAML at start of log.
- **Example schema:**

```
{
  "runId": "uuid",
  "timestamp": "2026-02-10T12:34:56Z",
  "cliVersion": "5.3.0",
  "model": "gpt-5.3-codex",
  "agentsHash": "sha256:abcd...",
  "rulesHash": "sha256:1234...",
  "commitSHA": "9f8e7d...",
  "role": "scaffold",
  "inputs": {
    "prompt": "...", "filesChanged": []
  },
  "outputs": { ... },
}
```

```
"metrics": { "tokensUsed": 2450 }  
}
```

Store this with the JSON events. SLSA provenance suggests capturing similar data: builder ID, buildType, parameters, resolvedDependencies, and outputs ¹⁶ ¹⁷. Here, treat Codex CLI as the builder: record `buildType: codex exec scaffold`, `builder.id: codex-cli-v5.3`, `parameters: {prompt:hashed, agentsHash, rulesHash}`, `subject: files output`.

Metrics & Observability

Record key metrics (as Prometheus counters or custom logs): - **Usage counts:** total Codex invocations, per agent role. - **Latency:** time per run. - **Tokens consumed** (Codex provides `/status` or count from logs). - **Outcome metrics:** number of PRs created vs updates, number of warnings/errors from Codex. - **Alerts:** if a Codex job fails or uses `danger-full-access` unexpectedly, trigger an alert to dev.

Send metrics to your monitoring (like Prometheus/Grafana). Example metric names: `codex_runs_total{role="scaffold"}`, `codex_tokens_consumed_total`.

6. Operational Patterns & Instance Roles

Define **instance roles** to compartmentalize responsibilities. Example roles:

- **Scaffold/Implementer:** Generates or updates code. Uses `workspace-write` sandbox. Prompt template: "Implement feature X" or scaffolding commands. Require approval (`on-change` or `always`).
- **Test-Fixer:** Focuses on tests and linting. Uses `read-only` or limited write (maybe `workspace-write` to add tests, but can `prompt` before changes). For example, runs after new code: "Write unit tests for changed files."
- **Security Auditor:** Scans code for vulnerabilities. `read-only`. Only asks for fixes (produces a report or suggestions as output).
- **Doc/ADR generator:** Updates docs or ADRs. `workspace-write` with on-request approval. Generates architecture decisions in `docs/`.

Each role has resource considerations. Limit simultaneous heavy roles (like two "refactor-all" tasks in parallel) because GitHub issue #10887 shows that parallel heavy Codex tasks can crash ³. Use job scheduling: e.g., queue heavy jobs sequentially.

Concurrency controls: In CI, you can use [concurrency groups](#) to prevent overlapping runs on the same target branch. In local, avoid literally simultaneous codex sessions in same directory; at least use separate CODEX_HOME to reduce collisions.

Rate/cost control: Monitor token consumption. If ChatGPT plan limits approached (see [20tL353-L361]), schedule long jobs in off-hours or use mini model. Pin model to Codex-Mini for small tasks to save budget ¹⁸.

7. CI/CD Integration (Examples)

Concrete GitHub Actions templates help cement these ideas.

GitHub Actions YAML: Example for two Codex jobs (adjusted for brevity):

```
name: Multi-Agent Codex CI
on: [pull_request]
permissions:
  contents: read
jobs:
  codex-scaffold:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - name: Setup Codex CLI
        uses: openai/setup-codex@v1
        with:
          version: 5.3.0
      - name: Run Scaffold Agent
        env:
          CODEX_HOME: ${{ runner.temp }}/codex_scaffold
          CODEX_API_KEY: ${{ secrets.CODEX_API_KEY }}
        run: |
          codex exec --json --profile scaffold --output-last-message scaffold.md \
            "Scaffold new endpoints for feature X in backend and frontend"
      - name: Upload Scaffold log
        uses: actions/upload-artifact@v3
        with:
          name: scaffold-log
          path: codex_scaffold/log.json
  codex-review:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - uses: openai/setup-codex@v1
        with: { version: 5.3.0 }
      - name: Run Reviewer Agent
        env:
          CODEX_HOME: ${{ runner.temp }}/codex_review
          CODEX_API_KEY: ${{ secrets.CODEX_API_KEY }}
        run: |
          codex exec --json --profile reviewer --output-last-message review.txt \
            "Review changes on this pull request and suggest fixes if needed"
      - name: Upload Review log
        uses: actions/upload-artifact@v3
        with:
          name: review-log
          path: codex_review/log.json
```

Key points: each job has its own CODEX_HOME, logs output, and does not auto-commit changes (no `--apply`). Instead, outputs are saved (we can use the CLI's `--output-last-message` to capture

suggestions in a file). A follow-up step could comment on the PR using this text (using GitHub Action or a small script), but **only after human review**.

Safe Auto-Fix Workflow: If you allow one agent to auto-fix (e.g. `codex exec --apply`), do so on a feature branch or fork, not on `main`. Better: have the agent open a PR itself with fixes, which a human then reviews. (We can script it: after running Codex with `--apply`, run `git push` to a new branch and use GitHub CLI to open a PR.)

GitLab CI: Similarly define parallel jobs with `CI matrix` and `artifacts`. The config syntax differs but concept is identical (use Docker or `openai/setup-codex` if available).

8. Local Dev Workflows

- **Project CODEX_HOME:** Encourage each developer to set `CODEX_HOME` inside project (e.g. `$PWD/.codex`) via `.env` file or launcher script. This isolates history per project.
- **Multiple CODEX_HOME:** You can manually switch CODEX_HOME in terminal or via scripts. Example shell script `run-codex-agent.sh`:

```
#!/bin/bash
CODEX_HOME="$PWD/.codex_$1" codex exec --profile $1 "$2"
```

- **Profiles:** Use `--profile` as in CI, to apply role-specific settings. Profiles are defined in `.codex/config.toml`.
- **Pre-commit Hooks:** Use `pre-commit` to run Codex checks (in *read-only* mode) before code commits. For example, a hook `codex-diff` could run `codex exec --no-apply --profile fixer "Check modified files for lint errors"` and fail on errors. This ensures issues caught early.
- **VS Code Integration:** The Codex VSCode extension can also use multiple agents (via worktrees/profiles), but keep settings consistent with CLI configs. Consider using Workspace settings to specify which skills or prompts the extension should use by default.

9. Testing & Review Strategy

- **Generated Code Tests:** Any code generated by Codex should have unit tests. E.g. if Codex adds a function, write tests in the same run or immediately after. Use `--output-schema` to have Codex generate JSON for API, then validate with schema tests.
- **Contract Testing:** For APIs, have OpenAPI schemas; use codex to generate initial spec, but verify with `openapi-generator` and real requests in CI.
- **Human Review:** Do **not** blindly commit Codex output. Always gate with a PR. Even if an agent suggests fixes, a developer must review diffs. Enforce this in policy.
- **Avoid Ephemeral Commits:** Disable auto-apply on mainline. If using `--apply`, do so in a side branch or local dev only.

10. Governance & Policy Checklist

- **Auto-commit vs PR:** Generally, restrict auto-commit. Only allow an agent to `--apply` changes in a draft branch. Most changes should require a PR with a human reviewer.

- **Code Ownership:** If agents generate code, clearly annotate generated sections (e.g. comments `// Generated by Codex CLI`). Assign ownership of each agent's output to a team member for accountability.
- **Secrets Rotation:** Rotate Codex API keys periodically (every few months) and immediately if a breach is suspected.
- **Artifact Retention:** Keep logs/artifacts for X months (depending on policy). Remove older ones to save space but keep enough for audits.
- **Incident Response:** If a Codex agent behaves maliciously (e.g. injected a bad patch), treat it like a supply-chain incident: revert commits, revoke tokens, check logs for intrusion. Codex tasks should be reproducible from logs (see next section).
- **Review Policies:** Require 2nd-party sign-off on changes from Codex if code impacts security/data. Maintain an ADR if adopting major agentic changes.
- **Compliance Note:** For PCI/GDPR, Codex has no PCI certification. If in-scope data is involved, do not send raw sensitive data to Codex prompts. Use placeholders or hashed tokens.

11. Migration & Rollback

If generated code causes regressions or security issues:

- **Revert:** Immediately revert the offending commit (e.g. `git revert`).
- **Investigate:** Check the Codex logs to identify prompts/inputs that led to it. Use the runId/commit to find JSON events. This is why metadata logging is crucial.
- **Reproduce:** To safely reproduce the run, pin down all inputs: same CLI version, model, AGENTS.md and rules version. Use JSONL and prompt as a *replay* (you may have to remove `--apply`).
- **Lockfiles:** For reproducibility, commit a lockfile for AGENTS.md or config (e.g. a script that regenerates them deterministically). Consider storing the entire prompt in a versioned spec.
- **Rollback Plan:** If large refactors, use the *Strangler Pattern*¹⁶: introduce new code around old code, switch over piece by piece. Codex can help generate adapters (anti-corruption layer) if needed.
- **Dual-write Pattern:** If splitting services in future, prepare adapters: run feature flags. Keep DB migrations backward-compatible while rewriting logic in new service.
- **Undo Generated Code:** If an agent made many file changes, and you must undo, use `codex exec --help` (hypothetically) or rely on `git revert`. No built-in "undo" yet – for now manual reverts.

12. Decision Criteria for Single vs Multiple Instances

Decide *when* to introduce multiple instances. Table 2 provides signals.

Table 2. When to stick with one agent vs add more agents.

Signals & Criteria	Continue Single Agent	Scale to Multiple Agents
Team size	1 developer (solo)	3+ developers (multiple domains)

Signals & Criteria	Continue Single Agent	Scale to Multiple Agents
Task complexity	Simple, single-domain tasks	Distinct domains (frontend/backend) 1
Workflow bottleneck	All tasks fit in one session/agent	Long-running runs, tasks queuing up 3
Code reviews backlog	Low (manual suffice)	High (introduce Codex agent for reviews)
Security requirements	Low (internal tools sufficient)	High (external agentic scans for code) 2
Budget for Codex	Very limited (avoid extra runs)	Adequate for parallel runs (monitor usage) 19
Maintainability risk	One set of agent instructions to update	Many roles to keep in sync (higher effort)

Moving from one to many is **high effort** (setting up orchestration, skills, rules) but can yield **high benefit** if workflow demands parallelism. Be cautious: issue [4] shows Codex stability can suffer under parallel load ³, so ensure you use separate processes and not one interactive session with background jobs.

13. Prioritised Next Steps

- Establish baseline:** Confirm a single Codex workflow (e.g. one agent scaffolding feature) works end-to-end with your tech stack.
- Set up AGENTS.md :** Draft project instructions (npm test rules, style guides, SQL lints) so all agents share context ⁸.
- Define roles & profiles:** Write a simple `.codex/config.toml` with profiles (`scaffold`, `fixer`, `reviewer`) and sandbox rules. (Use defaults from examples above.)
- Implement basic rules:** Create `.codex/rules/default.yaml` forbidding destructive commands (`rm -rf`, `sudo`) and requiring approval for `workspace-write` actions.
- CI pipeline skeleton:** Add a GitHub Actions job that runs `codex exec` (profile `scaffold`) with `--json` on PR. Validate it runs without errors.
- Logging:** Ensure the CI job archives the JSON log. Also make developers log CLI output locally (e.g. redirect to file).
- Secret management:** Store your Codex API key in GitHub Secrets. In local dev, use environment variable (`CODEX_API_KEY`).
- Pre-commit hook:** Integrate a pre-commit rule to run lint/ruff/tests, or even a read-only Codex lint check on changed files.
- Review process:** Draft a policy: all Codex changes require PRs. Train the team on how to review Codex-generated diffs.
- Monitor usage:** Keep an eye on Codex usage dashboard ²⁰, and set alerts for approaching limits. If needed, switch some tasks to GPT-5.1-Codex-Mini ¹⁸.

Following these steps builds a foundation. The full report below expands each area with detailed guidance, templates, and references.

References

- OpenAI Codex CLI docs (Configuration, Rules, AGENTS, Skills, Non-interactive, GitHub Action) [21](#)
[8](#) [22](#).
 - OpenAI blog: “*Unrolling the Codex agent loop*” [23](#).
 - GitHub Issues (Codex CLI): #3280 (multi-agent orchestration request) [1](#); #10887 (parallel tasks bug) [3](#).
 - Rezvani (2025): on multi-agent CI for security pipelines [2](#).
 - GitHub Actions Security (Least Privilege, Secrets) [12](#) [24](#).
 - OWASP Cheat Sheet: Secrets Management (least privilege, rotation) [25](#).
 - SLSA provenance spec (build metadata) [16](#) [17](#).
 - Trew Knowledge News (Feb 2026): multi-agent orchestration context [26](#).
 - OpenAI Codex Pricing & usage limits [27](#).
-

[1](#) Orchestration of multiple agents · Issue #3280 · openai/codex · GitHub

<https://github.com/openai/codex/issues/3280>

[2](#) Codex CLI - Complete Guide from Zero to Agents | Medium

<https://alirezarezvani.medium.com/openai-codex-cli-from-broken-install-to-multi-agent-orchestration-production-guide-07d8b7d513ef>

[3](#) Codex 5.3: Parallel Long-Running Tasks Can Cause Session Crashes / Truncated Output · Issue

#10887 · openai/codex · GitHub

<https://github.com/openai/codex/issues/10887>

[4](#) [8](#) Custom instructions with AGENTS.md

<https://developers.openai.com/codex/guides/agents-md/>

[5](#) [6](#) [7](#) [9](#) [21](#) Command line options

<https://developers.openai.com/codex/cli/reference/>

[10](#) [22](#) Agent Skills

<https://developers.openai.com/codex/skills/>

[11](#) [13](#) [14](#) [25](#) Secrets Management - OWASP Cheat Sheet Series

https://cheatsheetseries.owasp.org/cheatsheets/Secrets_Management_Cheat_Sheet.html

[12](#) [15](#) [24](#) Secure use reference - GitHub Docs

<https://docs.github.com/en/actions/reference/security/secure-use>

[16](#) [17](#) SLSA • Provenance

<https://slsa.dev/spec/v1.0/provenance>

[18](#) [19](#) [20](#) [27](#) Codex Pricing

<https://developers.openai.com/codex/pricing/>

[23](#) Unrolling the Codex agent loop | OpenAI

<https://openai.com/index/unrolling-the-codex-agent-loop/>

[26](#) AI This Week: Multi-Agent Orchestration Becomes Reality - Trew Knowledge Inc.

<https://trewknowledge.com/2026/02/06/ai-this-week-multi-agent-orchestration-becomes-reality/>